

Appendix A: Example **dsm** analysis

David L. Miller, M. Louise Burt, Eric A. Rexstad and Len Thomas

Preamble

This is a **knitr** document. The source for it contains everything you need to reproduce the analysis given here (aside from the data). The most recent version of this document can always be found at github.com/dill/mexico-data.

The data (which are included in the **dsm** package) consist of observations of dolphins in the Gulf of Mexico. The analysis is based on a dataset which is shipped with Distance 6.0. For convenience the data are bundled in an **R** friendly format, although all of the code necessary for creating the data from the Distance project files is available at the above URL.

The intention here is to highlight the features of the **dsm** package, rather than perform a full analysis of the data. For that reason, some important steps are not fully explored. Some familiarity with density surface modelling is assumed.

Before we start, we load the **dsm** package and set some options:

```
library(dsm)

## Loading required package: mgcv

## This is mgcv 1.7-22. For overview type 'help("mgcv-package")'.

## Loading required package: mrds

## Loading required package: optimx

## Loading required package: numDeriv

## Loading required package: Rsolnp

## Loading required package: truncnorm

## This is mrds 2.0.9 Built: R 2.15.2; ; 2012-12-10 10:15:50 UTC; unix
```

```
## Loading required package: ggplot2

## Loading required package: Distance

## This is dsm 2.0 Built: R 2.15.2; ; 2013-01-09 17:42:42 UTC; unix

# plotting options
gg.opts <- theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank(),
  panel.background = element_blank())
# make the results reproducible
set.seed(11123)
```

The data

Observation and segment data

All of the data for this analysis has been nicely pre-formatted and is shipped with dsm. Loading up that data, we can see that we have four data frames, the first few lines of each are shown:

```
data(mexdolphins)
attach(mexdolphins)
head(segdata)
```

##	latitude	longitude	Effort	Transect.Label	Sample.Label	depth	x
## 1	29.94	-86.93	13800	19960417	19960417-1	135.0	134159
## 2	29.84	-86.83	14000	19960417	19960417-2	147.7	143496
## 3	29.75	-86.74	14000	19960417	19960417-3	152.1	152050
## 4	29.66	-86.65	13900	19960417	19960417-4	163.8	161102
## 5	29.56	-86.57	13800	19960417	19960417-5	179.7	169553
## 6	29.49	-86.49	13800	19960417	19960417-6	188.5	176793

```
##      y
## 1 325561
## 2 314055
## 3 304324
## 4 293475
## 5 282984
## 6 275103

head(distdata)
```

##	object	size	distance	Effort	detected	beaufort	latitude	longitude
## 45	45	21	3296.6	36300	1	4	27.73	-86.00

```
## 61      61  150    929.2  17800      1      4    26.00   -87.63
## 63      63  125    6051.0  21000      1      2    26.01   -87.95
## 85      85   75    5499.7  21800      1      1    27.50   -90.45
## 114     114   50    7259.0  13400      1      3    27.41   -94.99
## 120     120   45    1454.8  20900      1      5    26.02   -95.97
##           x      y
## 45    228139  79258
## 61     69199 -113083
## 63     37046 -112197
## 85    -210016  54208
## 114   -658878  43337
## 120   -764824 -111005
```

```
head(obsdata)
```

```
##      object Sample.Label size distance Effort
## 45      45    19960421-9   21   3296.6  36300
## 61      61    19960423-7  150    929.2  17800
## 63      63    19960423-9  125   6051.0  21000
## 85      85    19960427-1   75   5499.7  21800
## 114     114    19960430-8   50   7259.0  13400
## 120     120    19960501-5   45   1454.8  20900
```

```
head(preddata)
```

```
##   latitude longitude depth      x      y width height
## 1    30.08    -87.58    35  70832 341079 32072  37065
## 2    30.08    -87.42    30  86868 341079 32072  37065
## 3    30.08    -87.25    27 102904 341079 32072  37065
## 4    30.08    -87.08    22 118940 341079 32072  37065
## 5    30.08    -86.92    46 134976 341079 32072  37065
## 6    29.92    -87.75    14  54888 322546 32126  37065
```

`distdata` holds the distance sampling data which will be used to fit the detection function. `segdata` holds the segment data: the transects have already been “chopped” into segments. `obsdata` holds the observations which have already been aggregated to the segments and `preddata` holds the prediction grid (which includes all the covariates that we need).

Typically it will be necessary to pre-allocate the observations to segments (as well as define the segment length) using a tool such as ArcGIS before starting an analysis using `dsm`.

The below figure shows the survey area with the transect lines overlaid (using data from `segdata`).

```

p <- qplot(data = survey.area, x = longitude, y = latitude, geom = "polygon",
           fill = I("lightblue"), ylab = "Latitude", xlab = "Longitude", alpha = I(0.7))
p <- p + coord_equal()
p <- p + gg.opts

p <- p + geom_line(aes(longitude, latitude, group = Transect.Label), data = segdata)

print(p)

```




Figure 1: The survey area with transect lines.

Converting units

It is important to ensure that the measurements to be used in the analysis are in compatible units, otherwise the resulting estimates will be incorrect or hard to interpret. Having all of our measurements in SI units from the outset removes the need for conversion later, making life much easier. All of the data is already in the appropriate units (Northings and Eastings: kilometres from ~ -88.32 longitude, ~ 27.02 latitude, which is the centroid of the study region, multiplied up by 1000 to get the result in metres for consistency).

It is also important to note that the area of the prediction cells must be included for prediction to work. Our data includes the width and height of each prediction

cell and we calculate the areas as needed. See below.

We give an example of converting the survey area here to show that this is a simple process:

```
# centroid
lon0 <- -88.31951
lat0 <- 27.01594

sa.tmp <- latlong2km(survey.area$longitude, survey.area$latitude, lon0 = lon0,
  lat0 = lat0)

survey.area <- data.frame(x = 1000 * sa.tmp$km.e, y = 1000 * sa.tmp$km.n)

rm(sa.tmp)
```

The function `latlong2km` makes this conversion simple (thanks to Simon N. Wood for providing code). The other data frames have already had their measurements appropriately converted. Note that by convention the directions are named `x` and `y`.

Using latitude and longitude when performing spatial smoothing can be problematic when certain smoother bases are used. In particular when bivariate isotropic bases are used the non-isotropic nature of latitude and longitude is inconsistent (moving one degree in one direction is not the same as moving one degree in the other).

Exploratory data analysis

Distance data

The top panels of the below figure show histograms of observed distances and group size and the bottom panels show the relationship between observed distance and observed group size, and the relationship between observed distance and Beaufort sea state. The plots show that there is some relationship between size and observed distance to be explored since as one would expect, fewer smaller groups seem to be seen at larger distances.

```
# save graphics options
o <- par("mfrow")
par(mfrow = c(2, 2))

# histograms
hist(distdata$distance, main = "", xlab = "Distance (m)")
hist(distdata$size, main = "", xlab = "Group size")
```

```

# plots of distance vs. size
plot(distdata$distance, distdata$size, main = "", xlab = "Distance (m)", ylab = "Group size",
      pch = 19, cex = 0.5, col = rgb(0.74, 0.74, 0.74, 0.7))

# lm fit
l.dat <- data.frame(distance = seq(0, 8000, len = 1000))
lo <- lm(size ~ distance, data = distdata)
lines(l.dat$distance, as.vector(predict(lo, l.dat)))

plot(distdata$distance, distdata$beaufort, main = "", xlab = "Distance (m)",
      ylab = "Beaufort sea state", pch = 19, cex = 0.5, col = rgb(0.74, 0.74,
        0.74, 0.7))

# restore graphics options
par(o)

## NULL

```

Spatial data

Looking separately at the spatial data without thinking about the distances, we can see the distribution of group size in space in the below figure. Circle size indicates the size of the group in the observation. There are rather large areas with no observations, which might cause our variance estimates to be rather large.

```

p <- qplot(data = survey.area, x = x, y = y, geom = "polygon", ylab = "y", xlab = "x",
           alpha = I(0.7), fill = I("lightblue"))
p <- p + gg.opts
p <- p + coord_equal()
p <- p + labs(size = "Group size")
p <- p + geom_line(aes(x, y, group = Transect.Label), data = segdata)
p <- p + geom_point(aes(x, y, size = size), data = distdata, colour = "red",
                   alpha = I(0.7))
print(p)

```

We will use depth later as an explanatory covariate in our spatial model. The next plot shows the raw depth data.

```

p <- ggplot(preddata)
p <- p + gg.opts

```




Figure 2: Exploratory plot of the distance sampling data. Top row, left to right: histograms of distance and group size; bottom row: plot of distance against group size and plot of distances against Beaufort sea state.




Figure 3: The survey area. Point size is proportional to the group size for each observation.


```

p <- p + coord_equal()
p <- p + labs(fill = "Depth", x = "x", y = "y")
p <- p + geom_tile(aes(x = x, y = y, fill = depth, width = width, height = height))
print(p)

```




Figure 4: Plot of depth values over the survey area.

Combining the above two plots yields:

```

p <- ggplot(preddata)
p <- p + gg.opts
p <- p + coord_equal()
p <- p + labs(fill = "Depth", x = "x", y = "y", size = "Group size")
p <- p + geom_tile(aes(x = x, y = y, fill = depth, width = width, height = height))
p <- p + geom_line(aes(x, y, group = Transect.Label), data = segdata)
p <- p + geom_point(aes(x, y, size = size), data = distdata, colour = "red",
  alpha = I(0.7))
print(p)

```

This plot shows that we don't seem to have many observations in the very shallow areas near the shore. This should make us skeptical of predictions in those areas.




Figure 5: Plot of depth values over the survey area with transects and observations overlaid.

Estimating the detection function

We use the `ds` function in `Distance` to fit the detection function. First, loading the `Distance` library:

```
library(Distance)
```

We can then fit a detection function with hazard-rate key with no adjustment terms:

```
hr.model <- ds(distdata, max(distdata$distance), key = "hr", adjustment = NULL)
```

```
## Fitting hazard-rate key function
```

```
## AIC= 841.253
```

```
## No survey area information supplied, only estimating detection function.
```

```
summary(hr.model)
```

```
##
## Summary for distance analysis
## Number of observations : 47
## Distance range       : 0 - 7847
##
## Model : Hazard-rate key function
## AIC   : 841.3
##
## Detection function parameters
## Scale Coefficients:
##           estimate      se
## (Intercept)  7.983 0.9532
##
## Shape parameters:
##           estimate      se
## (Intercept)    0 0.7835
##
##           Estimate      SE      CV
## Average p      0.5913 0.2224 0.3762
## N in covered region 79.4879 30.8073 0.3876
```

```
plot(hr.model)
```

For brevity, detection function model selection has been omitted here. In practise we would fit many different forms for the detection function, below we fit a detection function with size as a covariate, but for now we stick to a simple model.




Figure 6: Plot of the fitted detection function.

Fitting a DSM

Before fitting a `dsm` model, the data must be segmented; this consists of chopping up the transects and attributing counts to each of the segments. As mentioned above, these data have already been segmented.

A simple model

We begin with a very simple model. We assume that the number of individuals in each segment are quasi-Poisson distributed and that they are a smooth function of their spatial coordinates (note that the formula is exactly as one would specify to `gam` in `mgcv`). The abundance of groups rather than individuals can be estimated by setting `group=TRUE` (though we ignore this here).

Running the model:

```
mod1 <- dsm(N ~ s(x, y), hr.model$ddf, segdata, obsdata)
summary(mod1)

##
## Family: quasipoisson
```

```
## Link function: log
##
## Formula:
## N ~ s(x, y) + offset(off.set)
## <environment: 0x1075eabf0>
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -18.409      0.394   -46.7   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df    F p-value
## s(x,y) 26.1    28.2 4.42 4.1e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.113   Deviance explained =   44%
## GCV score = 42.985   Scale est. = 37.611     n = 387
```

We can then make the predictions over the the grid and calculate abundance. First we must create the offset (the area of each grid cell, which is 444km²).

```
off.set <- 444 * 1000 * 1000
mod1.pred <- predict(mod1, preddata, off.set)
```

Below is a map of the predicted abundance. Before plotting, we bind on the predictions to the data used to create them:

```
pp <- cbind(preddata, mod1.pred)
```

Then plotting:

```
p <- ggplot(pp) + gg.opts
p <- p + geom_tile(aes(x = x, y = y, fill = mod1.pred, width = width, height = height))
p <- p + coord_equal()
p <- p + scale_fill_gradientn(colours = heat_hcl(1000))

## Error: could not find function "heat_hcl"

p <- p + geom_path(aes(x = x, y = y), data = survey.area)
p <- p + labs(fill = "Abundance")
print(p)
```




Figure 7: Predicted density surface for mod1.

We can calculate abundance over the survey area by simply summing these predictions:

```
sum(mod1.pred)
```

```
## [1] 47034
```

We can also look at diagnostic plots for the model.

```
gam.check(mod1)
```

```
##
## Method: GCV   Optimizer: outer newton
## full convergence after 4 iterations.
## Gradient range [4.281e-07,4.281e-07]
## (score 42.99 & scale 37.61).
## Hessian positive definite, eigenvalue range [0.4516,0.4516].
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
```

```
##
##           k'   edf k-index p-value
## s(x,y) 29.00 26.06   1.02   0.98
```

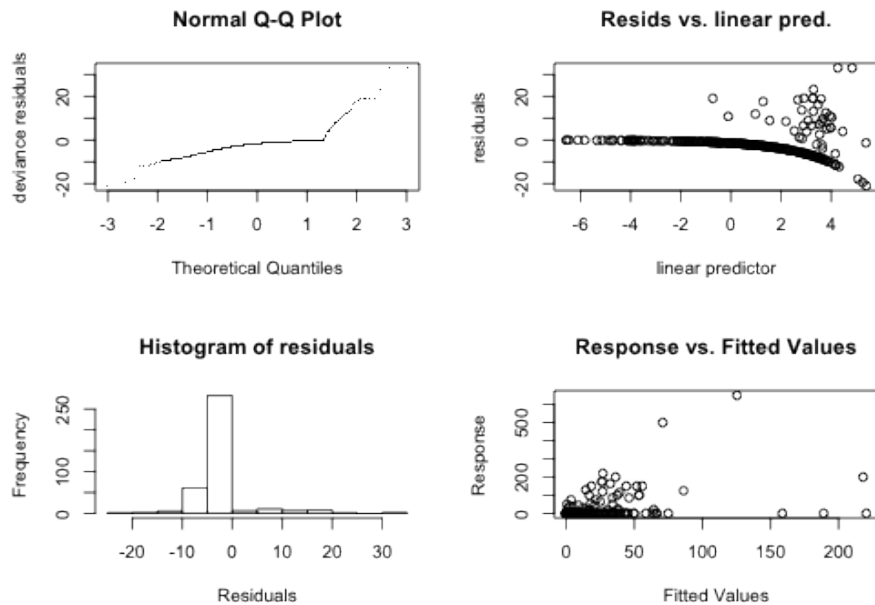


Figure 8: Diagnostic plots for mod1.

These show that there is some deviation in the Q-Q plot. The “line” of points in the plot of the residuals vs. linear predictor plot corresponds to the zeros in the data.

We can use the approach of Williams et al (2011), which accounts for uncertainty in detection function estimation.

```
preddata.varprop <- split(preddata, 1:nrow(preddata))
offset.varprop <- as.list(rep(off.set, nrow(preddata)))
mod1.varprop <- dsm.var.prop(mod1, pred.data = preddata.varprop, off.set = offset.varprop)
```

Calling `summary` will give some information about uncertainty estimation:

```
summary(mod1.varprop)
```

```
## Summary of uncertainty in a density surface model calculated
## by variance propagation.
```

```
##
## Quantiles of differences between fitted model and variance model
##      Min.    1st Qu.      Median        Mean     3rd Qu.      Max.
## -4.69e-12 -1.57e-13 -1.40e-14 -2.11e-13  4.00e-15  2.70e-13
##
## Approximate asymptotic confidence interval:
##      5% Mean  95%
## 29962 47034 73832
## (Using delta method)
##
##
## Point estimate           : 47034
## Standard error           : 10966
## Coefficient of variation   : 0.2331
```

We can also make a plot of the CVs:

```
plot(mod1.varprop, xlab = "Easting", ylab = "Northing")
```




Figure 9: Plot of the coefficient of variation for the study area.

Adding another covariate to the spatial model

The data set also contains a `depth` covariate (which we plotted above). We can include in the model very simply:

```
mod2 <- dsm(N ~ s(x, y, k = 10) + s(depth, k = 20), hr.model, segdata, obsdata,
            select = TRUE)
summary(mod2)

##
## Family: quasipoisson
## Link function: log
##
## Formula:
## N ~ s(x, y, k = 10) + s(depth, k = 20) + offset(off.set)
## <environment: 0x108e3df08>
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -18.823      0.734   -25.6   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##               edf Ref.df    F p-value
## s(x,y)         3.88     9 2.06 0.00018 ***
## s(depth)      10.52    19 3.15 9.9e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.0929   Deviance explained =   34%
## GCV score =  46.27   Scale est. = 42.967    n = 387
```

By setting the `k` parameter we specify the largest complexity for that smooth term in the model; as long as this is high enough (i.e. the number in the `edf` column is not too close to that in the `Ref.df` column, we can be sure that there is enough flexibility. However, it may sometimes be necessary to set `k` to be lower than this to limit the influence of (for example) spatial smoothers in the model.

Setting `select=TRUE` here imposes extra shrinkage terms on each smooth in the model (allowing smooth terms to be removed from the model during fitting, see `?gam` for more information). Although this is not particularly useful here, this can be a good way (along with looking at p-values) to perform term selection.

Again we can plot the predictions from this model:

```

mod2.pred <- predict(mod2, preddata, off.set)
pp <- cbind(preddata, mod2.pred)
p <- ggplot(pp) + gg.opts
p <- p + labs(fill = "Abundance")
p <- p + geom_tile(aes(x = x, y = y, fill = mod2.pred, width = width, height = height))
p <- p + coord_equal()
p <- p + scale_fill_gradientn(colours = heat_hcl(20))

## Error: could not find function "heat_hcl"

p <- p + geom_path(aes(x = x, y = y), data = survey.area)
print(p)

```




Figure 10: Predicted density surface for mod2.

We can also look at the relationship between depth and abundance:

```
plot(mod2, select = 2)
```

Omitting the argument `select` gives plots of all the smooth terms, one at a time.




Figure 11: Plot of the smooth of depth in mod2.

A more complicated model

Tweedie

Response distributions other than the quasi-Poisson can be used, for example the Tweedie distribution. If the Tweedie is used, then the `p` parameter must be specified. Mark Bravington (via personal communication) suggests the use of `p=1.2` for marine mammal work. The choice of `p` is only sensitive to the first decimal place, so a quick search can be performed by simply comparing the score of the resulting models.

```
mod3 <- dsm(N ~ s(x, y), hr.model, segdata, obsdata, family = Tweedie(p = 1.2))
summary(mod3)

##
## Family: Tweedie(1.2)
## Link function: log
##
## Formula:
## N ~ s(x, y) + offset(off.set)
## <environment: 0x10a1d81c8>
##
```

```
## Parametric coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -18.981      0.405   -46.9   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##           edf Ref.df    F p-value
## s(x,y) 27.3   28.7 6.48  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.0531   Deviance explained = 47.9%
## GCV score = 23.284   Scale est. = 20.239    n = 387
```

Note that we have to fully specify the family just as in `mgcv`.

As well as looking at the GCV/UBRE/REML score of the model to assess the value of p we can also look at a plot of the square root of the absolute value of the residuals versus the fitted values gives the following plot:

```
plot(sqrt(abs(residuals(mod3))), predict(mod3), xlab = "Square root of the absolute value of
      ylab = "Fitted values")
```

where as setting $p=1.7$:

```
mod3 <- dsm(N ~ s(x, y), hr.model, segdata, obsdata, family = Tweedie(p = 1.7))
summary(mod3)

##
## Family: Tweedie(1.7)
## Link function: log
##
## Formula:
## N ~ s(x, y) + offset(off.set)
## <environment: 0x11ed2a460>
##
## Parametric coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -24.267      0.671   -36.2   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##           edf Ref.df    F p-value
```



Figure 12: Plot of absolute value of the residuals versus the fitted values for the Tweedie model when $p=1.2$

```
## s(x,y) 28.8      29 14.3 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  -9.81   Deviance explained = 53.4%
## GCV score = 9.3865   Scale est. = 8.0932     n = 387
```

we obtain the following plot, which appears to be much flatter:

```
plot(sqrt(abs(residuals(mod3))), predict(mod3), xlab = "Square root of the absolute value of
      ylab = "Fitted values")
```

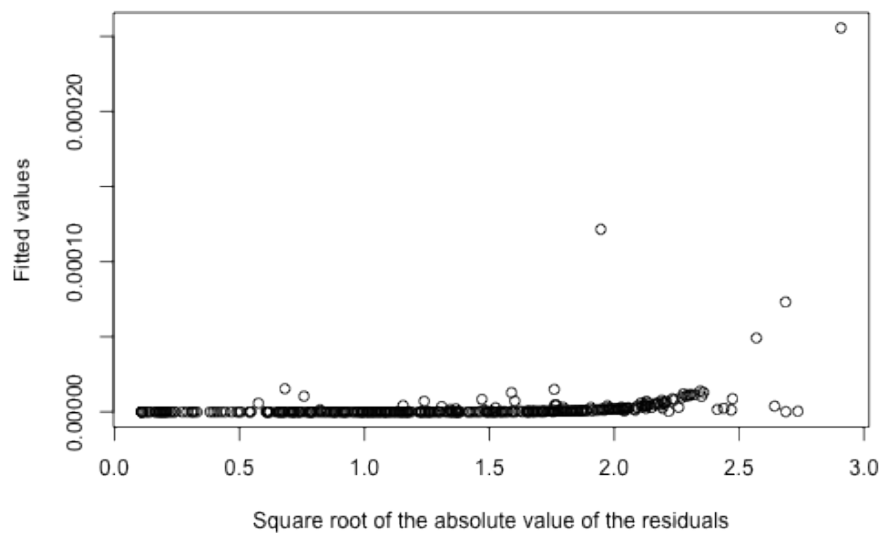


Figure 13: Plot of absolute value of the residuals versus the fitted values for the Tweedie model when $p=1.7$, note that this plot is much flatter than the previous plot.

Note also the improvement in GCV score.

In general, a “good” value can be found by simply plotting the above along with the GCV score for the model for values of p between 1.1 and 1.9 and looking for the best GCV and “flattest” plot.

Soap film smoothing

To account for a complex region (e.g. a region that includes peninsulae) we can use the soap film smoother. This is provided by the soap film smoother (see references, below).

In order to use a soap film smoother for the spatial part of the model we must create a set of knots for the smoother to use. This is easily done using the `make_soap_grid()` function in `dsm`:

```
soap.knots <- make.soapgrid(survey.area, c(11, 6))

## Warning: the condition has length > 1 and only the first element will be
## used

# knot 11 is not outside but is too close according to soap...
soap.knots <- soap.knots[-11, ]
```

where the second argument specifies the size of the grid that will be used to create the knots (knots in the grid outside of `survey.area` are removed).

As we saw in the exploratory analysis, some of the transect lines are outside of the survey area. These will cause the soap film smoother to fail, so we remove them:

```
x <- segdata$x
y <- segdata$y
onoff <- inSide(x = x, y = y, bnd = survey.area)

## Warning: the condition has length > 1 and only the first element will be
## used

rm(x, y)
segdata <- segdata[onoff, ]
```

We can run a model with both the depth covariate along with a spatial (soap film) smooth:

```
mod4 <- dsm(N ~ s(x, y, bs = "so", k = 10, xt = list(bnd = list(survey.area))) +
  s(depth), hr.model, segdata, obsdata, knots = soap.knots)

## Loading required package: Matrix

## Loading required package: lattice

## Attaching package: 'Matrix'
```

```

## The following object(s) are masked from 'package:stats':
##
## toeplitz

summary(mod4)

##
## Family: quasipoisson
## Link function: log
##
## Formula:
## N ~ s(x, y, bs = "so", k = 10, xt = list(bnd = list(survey.area))) +
##       s(depth) + offset(off.set)
## <environment: 0x10a2be6d0>
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -19.40      0.62    -31.3   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##               edf Ref.df    F p-value
## s(x,y)      23.05  29.00 2.63  2e-07 ***
## s(depth)    7.08   7.89 2.21  0.027 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.183   Deviance explained = 45.9%
## GCV score = 45.227   Scale est. = 38.341    n = 365

```

Not that the `k` argument now refers to the complexity of the boundary smooth in the soap film and the complexity of the film is controlled by the knots given in the `xt` argument.

Comparing predictions from the model that included a smooth of depth, we can see that the soap film has prevented some of the extreme values (especially in the lower right corner of the survey area).

```

mod4.pred <- predict(mod4, preddata, off.set)
pp <- cbind(preddata, mod4.pred)

p <- ggplot(pp) + gg.opts
p <- p + geom_tile(aes(x = x, y = y, fill = mod4.pred, width = width, height = height))
p <- p + coord_equal()
p <- p + scale_fill_gradientn(colours = heat_hcl(20))

```



```
## Error: could not find function "heat_hcl"

p <- p + geom_path(aes(x = x, y = y), data = survey.area)
p <- p + labs(fill = "Abundance")
print(p)
```




Figure 14: Predicted density surface for mod4.

Adding covariates to the detection function

It is common to include covariates in the detection function (so-called Multiple Covariate Distance Sampling or MCDS). In this dataset there are two covariates which are collected on each individual: Beaufort sea state and size. For brevity we fit only a hazard-rate detection functions with the sea state included as a factor covariate as follows:

```
hr.beau.model <- ds(distdata, max(distdata$distance), formula = ~as.factor(beaufort),
  key = "hr", adjustment = NULL)
summary(hr.beau.model)

##
```

```
## Summary for distance analysis
## Number of observations : 47
## Distance range      : 0 - 7847
##
## Model : Hazard-rate key function
## AIC   : 843.7
##
## Detection function parameters
## Scale Coefficients:
##              estimate      se
## (Intercept)    7.66318  1.077
## as.factor(beaufort)2  2.27968 17.367
## as.factor(beaufort)3  0.28606  1.019
## as.factor(beaufort)4  0.07174  1.223
## as.factor(beaufort)5 -0.36399  1.537
##
## Shape parameters:
##              estimate      se
## (Intercept)    0.3004  0.518
##
##              Estimate      SE      CV
## Average p          0.5421  0.1751 0.3229
## N in covered region 86.6957 29.4133 0.3393
```

Note that we opt to not use adjustments with the covariate model since we cannot ensure that the detection function will remain monotonic if both covariates and adjustment terms are used.

We fit the model as above, simply replacing the detection function model object and changing the response to be `Nhat` so the abundances are estimated per segment:

```
mod5 <- dsm(Nhat ~ s(x, y), hr.beau.model, segdata, obsdata)
summary(mod5)
```

```
##
## Family: quasipoisson
## Link function: log
##
## Formula:
## Nhat ~ s(x, y) + offset(off.set)
## <environment: 0x123e7fd58>
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -17.919      0.274   -65.5   <2e-16 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##           edf Ref.df   F p-value
## s(x,y) 25.2   27.8 3.6 8.9e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.119   Deviance explained = 40.7%
## GCV score = 81.573   Scale est. = 71.124     n = 365
```

Note that for models where there are covariates at the individual level we cannot calculate the variance via the variance propagation method of Williams et al (2011), but rather must resort to the moving block bootstrap. Other than this, all of the above models can be fitted.

A plot of predictions from the covariate model:

```
mod5.pred <- predict(mod5, preddata, off.set)
pp <- cbind(preddata, mod5.pred)

p <- ggplot(pp) + gg.opts
p <- p + geom_tile(aes(x = x, y = y, fill = mod5.pred, width = width, height = height))
p <- p + coord_equal()
p <- p + scale_fill_gradientn(colours = heat_hcl(20))

## Error: could not find function "heat_hcl"

p <- p + geom_path(aes(x = x, y = y), data = survey.area)
p <- p + labs(fill = "Abundance")
print(p)
```

Conclusions

This document has hopefully shown that the **dsm** package is a versatile and relatively easy-to-use package for the analysis of spatial distance sampling data. Note that there are many possible models that can be fitted using **dsm** and that the aim here was to show just a few of the options. Results from the models can be rather different, so care must be taken in performing model selection and discrimination.




Figure 15: Predicted density surface for mod5.

Notes

- `Distance` is available at <http://github.com/dill/Distance> as well as on CRAN.
- `dsm` is available (along with some documentation and hints) at <http://github.com/dill/dsm>, as well as on CRAN.

References

- Williams, R., Hedley, S.L., Branch, T.A., Bravington, M.V., Zerbini, A.N. & Findlay, K.P. (2011) Chilean Blue Whales as a Case Study to Illustrate Methods to Estimate Abundance and Evaluate Conservation Status of Rare Species. *Conservation Biology*, 25, 526–535.
- Hedley, S.L. & Buckland, S.T. (2004) Spatial models for line transect sampling. *Journal of Agricultural, Biological, and Environmental Statistics*, 9, 181–199.